

Hierarchical Clustering

Dutt Thakkar

2023-05-15

Loading the dataset

```
data=read.csv("/Users/duttthakkar/Desktop/Cereals.csv")
head(data)
```

```
##              name mfr type calories protein fat sodium fiber carbo
## 1      100%_Bran   N    C      70       4   1   130   10.0   5.0
## 2    100%_Natural_Bran Q    C     120      3   5    15    2.0   8.0
## 3          All-Bran   K    C      70       4   1   260    9.0   7.0
## 4 All-Bran_with_Extra_Fiber K    C      50      4   0   140   14.0   8.0
## 5        Almond_Delight R    C     110      2   2   200    1.0  14.0
## 6  Apple_Cinnamon_Cheerios G    C     110      2   2   180    1.5  10.5
##  sugars potass vitamins shelf weight cups  rating
## 1      6    280      25     3      1 0.33 68.40297
## 2      8    135       0     3      1 1.00 33.98368
## 3      5    320      25     3      1 0.33 59.42551
## 4      0    330      25     3      1 0.50 93.70491
## 5      8     NA      25     3      1 0.75 34.38484
## 6     10     70      25     1      1 0.75 29.50954
```

Viewing the summary and structure of that dataset

```
summary(data)
```

```
##      name          mfr          type          calories
## Length:77      Length:77      Length:77      Min.   : 50.0
## Class :character Class :character Class :character 1st Qu.:100.0
## Mode  :character Mode  :character Mode  :character Median :110.0
##                                     Mean   :106.9
##                                     3rd Qu.:110.0
##                                     Max.   :160.0
##
##      protein      fat      sodium      fiber
## Min.   :1.000    Min.   :0.000    Min.   : 0.0    Min.   : 0.000
## 1st Qu.:2.000    1st Qu.:0.000    1st Qu.:130.0   1st Qu.: 1.000
## Median :3.000    Median :1.000    Median :180.0   Median : 2.000
## Mean   :2.545    Mean   :1.013    Mean   :159.7   Mean   : 2.152
## 3rd Qu.:3.000    3rd Qu.:2.000    3rd Qu.:210.0   3rd Qu.: 3.000
## Max.   :6.000    Max.   :5.000    Max.   :320.0   Max.   :14.000
##
##      carbo      sugars      potass      vitamins
## Min.   : 5.0    Min.   : 0.000    Min.   : 15.00   Min.   : 0.00
## 1st Qu.:12.0    1st Qu.: 3.000    1st Qu.: 42.50   1st Qu.: 25.00
## Median :14.5    Median : 7.000    Median : 90.00   Median : 25.00
## Mean   :14.8    Mean   : 7.026    Mean   : 98.67   Mean   : 28.25
## 3rd Qu.:17.0    3rd Qu.:11.000    3rd Qu.:120.00   3rd Qu.: 25.00
## Max.   :23.0    Max.   :15.000    Max.   :330.00   Max.   :100.00
## NA's   :1      NA's   :1      NA's   :2
##      shelf      weight      cups      rating
## Min.   :1.000    Min.   :0.50    Min.   :0.250    Min.   :18.04
## 1st Qu.:1.000    1st Qu.:1.00    1st Qu.:0.670    1st Qu.:33.17
## Median :2.000    Median :1.00    Median :0.750    Median :40.40
## Mean   :2.208    Mean   :1.03    Mean   :0.821    Mean   :42.67
## 3rd Qu.:3.000    3rd Qu.:1.00    3rd Qu.:1.000    3rd Qu.:50.83
## Max.   :3.000    Max.   :1.50    Max.   :1.500    Max.   :93.70
##
```

```
str(data)
```

```
## 'data.frame': 77 obs. of 16 variables:
## $ name : chr "100%_Bran" "100%_Natural_Bran" "All-Bran" "All-Bran_with_Extra_Fiber" ...
## $ mfr : chr "N" "Q" "K" "K" ...
## $ type : chr "C" "C" "C" "C" ...
## $ calories: int 70 120 70 50 110 110 110 130 90 90 ...
## $ protein : int 4 3 4 4 2 2 2 3 2 3 ...
## $ fat : int 1 5 1 0 2 2 0 2 1 0 ...
## $ sodium : int 130 15 260 140 200 180 125 210 200 210 ...
## $ fiber : num 10 2 9 14 1 1.5 1 2 4 5 ...
## $ carbo : num 5 8 7 8 14 10.5 11 18 15 13 ...
## $ sugars : int 6 8 5 0 8 10 14 8 6 5 ...
## $ potass : int 280 135 320 330 NA 70 30 100 125 190 ...
## $ vitamins: int 25 0 25 25 25 25 25 25 25 ...
## $ shelf : int 3 3 3 3 3 1 2 3 1 3 ...
## $ weight : num 1 1 1 1 1 1 1 1.33 1 1 ...
## $ cups : num 0.33 1 0.33 0.5 0.75 0.75 1 0.75 0.67 0.67 ...
## $ rating : num 68.4 34 59.4 93.7 34.4 ...
```

Loading required packages

```
library(caret)
```

```
## Loading required package: ggplot2
```

```
## Loading required package: lattice
```

```
library(corrplot)
```

```
## corrplot 0.92 loaded
```

```
library(ggcorrplot)
```

```
library(tidyverse)
```

```
## — Attaching core tidyverse packages ————— tidyverse 2.0.0 —
## ✓ dplyr      1.1.0      ✓ readr      2.1.4
## ✓ forcats    1.0.0      ✓ stringr    1.5.0
## ✓ lubridate  1.9.2      ✓ tibble     3.1.8
## ✓ purrr      1.0.1      ✓ tidyr      1.3.0
```

```
## — Conflicts ————— tidyverse_conflicts() —  
## * dplyr::filter() masks stats::filter()  
## * dplyr::lag()     masks stats::lag()  
## * purrr::lift()    masks caret::lift()  
## i Use the [8;;http://conflicted.r-lib.org/conflicted package]8;; to force all conf  
licts to become errors
```

```
library(tidyr)  
library(dplyr)  
library(e1071)  
library(reshape2)
```

```
##  
## Attaching package: 'reshape2'  
##  
## The following object is masked from 'package:tidyr':  
##  
##     smiths
```

```
library(factoextra)
```

```
## Welcome! Want to learn more? See two factoextra-related books at https://goo.gl/ve3WBa
```

```
library(cluster)  
library(cowplot)
```

```
##  
## Attaching package: 'cowplot'  
##  
## The following object is masked from 'package:lubridate':  
##  
##     stamp
```

```
library(pander)  
library(kernlab)
```

```
##  
## Attaching package: 'kernlab'  
##  
## The following object is masked from 'package:purrr':  
##  
##     cross  
##  
## The following object is masked from 'package:ggplot2':  
##  
##     alpha
```

```
library(FactoMineR)
```

Data Preprocessing. Remove all cereals with missing values.

```
dim(data)
```

```
## [1] 77 16
```

```
c_d2=na.omit(data)  
dim(c_d2)
```

```
## [1] 74 16
```

There were 4 missing values in the dataset Assigning row names to the cereal column

```
c_d3 = as.data.frame(c_d2)  
row.names(c_d3) = c_d3[,1]  
c_d4 = c_d3[,-1]
```

Only selecting numerical values and removing categorical variables

```
c_d5 = c_d4[, c(3:11,13:15)]
```

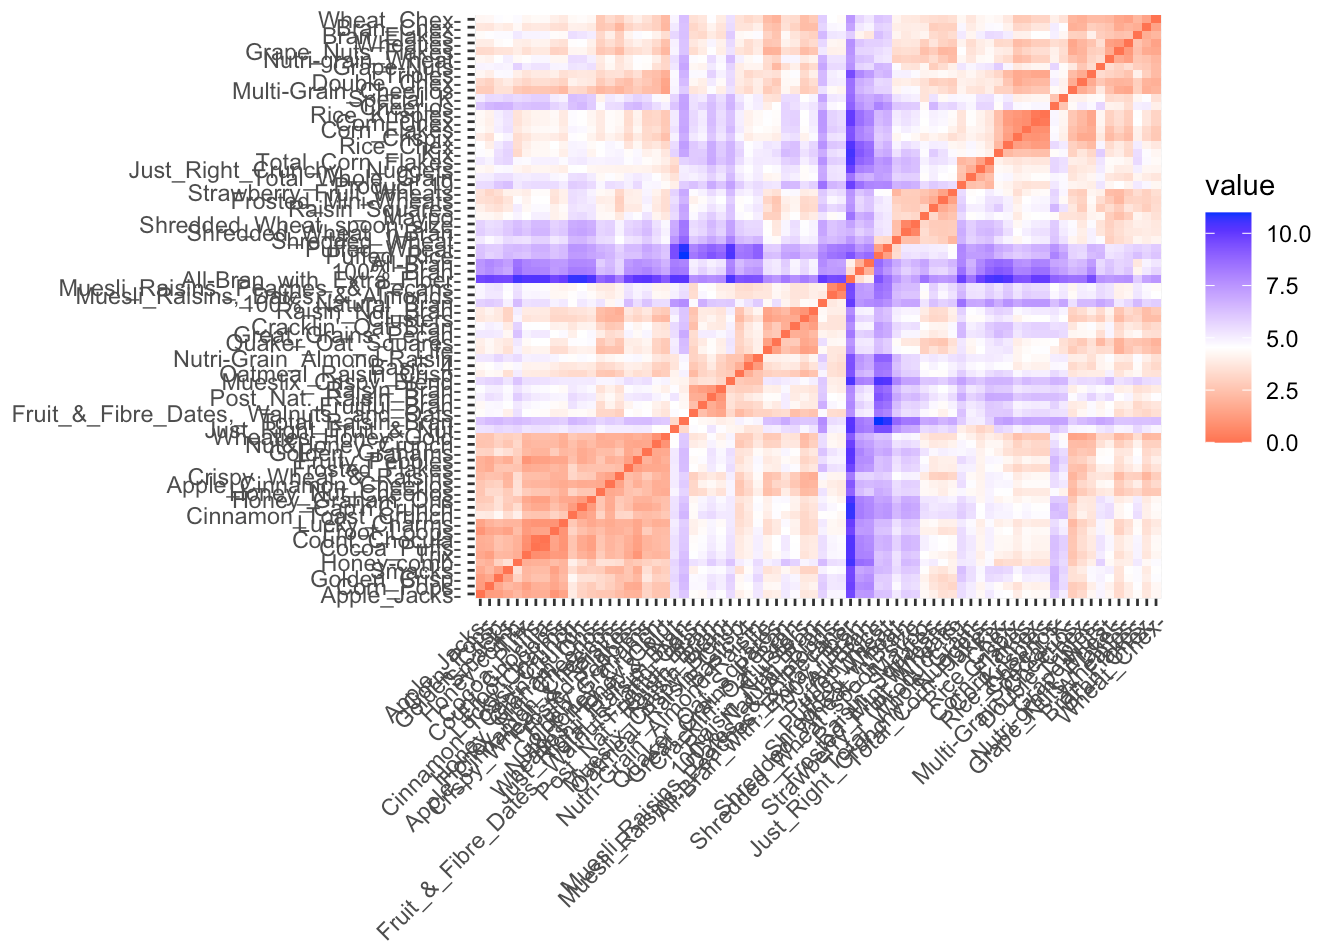
Normalizing the data using the scale function

```
c_d5 = scale(c_d5)
head(c_d5)
```

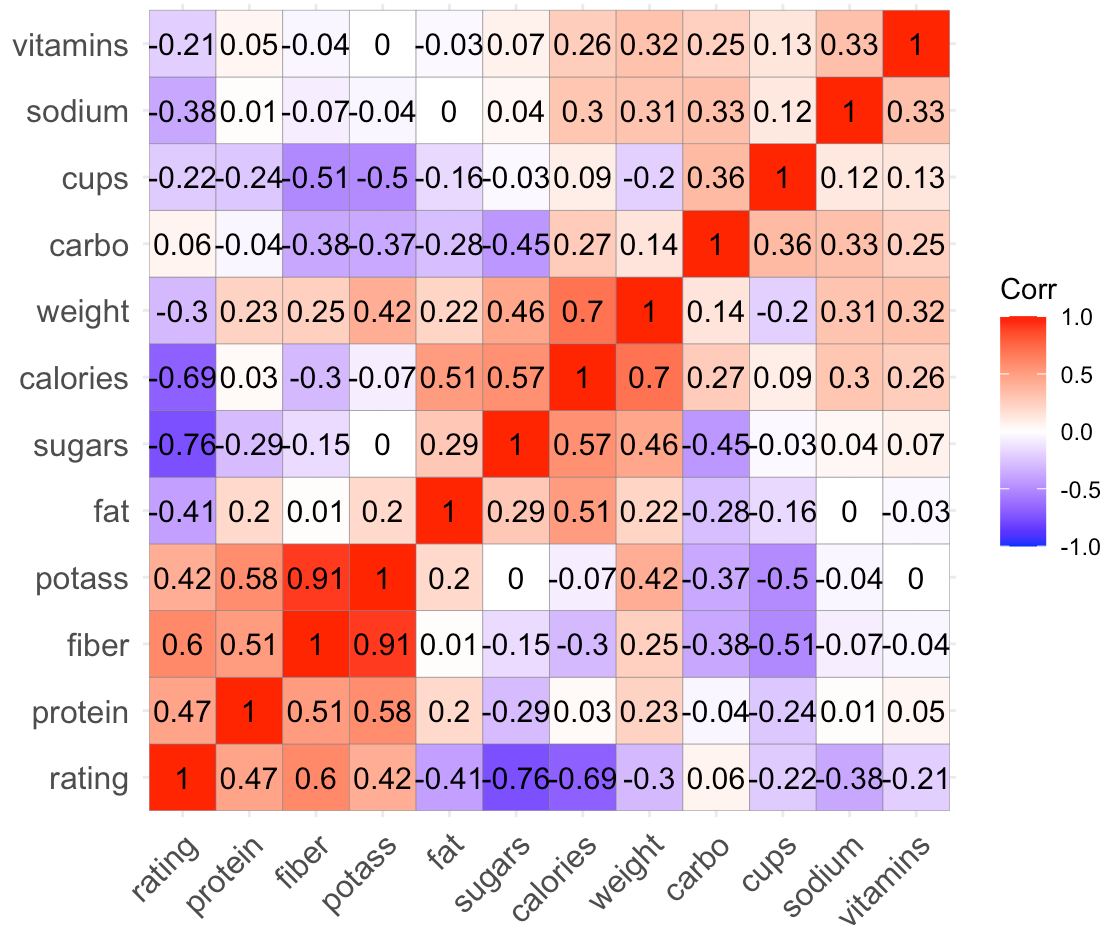
```
##              calories    protein      fat    sodium
## 100%_Bran      -1.8659155  1.3817478  0.0000000 -0.3910227
## 100%_Natural_Bran  0.6537514  0.4522084  3.9728810 -1.7804186
## All-Bran      -1.8659155  1.3817478  0.0000000  1.1795987
## All-Bran_with_Extra_Fiber -2.8737823  1.3817478 -0.9932203 -0.2702057
## Apple_Cinnamon_Cheerios  0.1498180 -0.4773310  0.9932203  0.2130625
## Apple_Jacks    0.1498180 -0.4773310 -0.9932203 -0.4514312
##              fiber      carbo      sugars    potass
## 100%_Bran      3.22866747 -2.5001396 -0.2542051  2.5605229
## 100%_Natural_Bran -0.07249167 -1.7292632  0.2046041  0.5147738
## All-Bran      2.81602258 -1.9862220 -0.4836096  3.1248675
## All-Bran_with_Extra_Fiber 4.87924705 -1.7292632 -1.6306324  3.2659536
## Apple_Cinnamon_Cheerios -0.27881412 -1.0868662  0.6634132 -0.4022862
## Apple_Jacks    -0.48513656 -0.9583868  1.5810314 -0.9666308
##              vitamins    weight      cups    rating
## 100%_Bran      -0.1818422 -0.2008324 -2.0856582  1.8549038
## 100%_Natural_Bran -1.3032024 -0.2008324  0.7567534 -0.5977113
## All-Bran      -0.1818422 -0.2008324 -2.0856582  1.2151965
## All-Bran_with_Extra_Fiber -0.1818422 -0.2008324 -1.3644493  3.6578436
## Apple_Cinnamon_Cheerios -0.1818422 -0.2008324 -0.3038480 -0.9165248
## Apple_Jacks    -0.1818422 -0.2008324  0.7567534 -0.6553998
```

Question 1 (part A): Apply hierarchical clustering to the data using Euclidean distance to the normalized measurements and looking at the correlation values by plotting the corrplot

```
distance_table <- get_dist(c_d5)
fviz_dist(distance_table)
```



```
corr_plot = cor(c_d5)
ggcorrplot(corr_plot, outline.color = "grey50", lab = TRUE, hc.order = TRUE, type = "full")
```



#Sugar and calories are highly negatively correlated with rating. Also, Potass is highly positively correlated with fiber and Protein.

Question 1 (part B): comparing hierarchical clustering with different linkages: single, average, complete and ward.

```
# Hierarchical clustering using Complete Linkage
hc1 <- hclust(distance_table, method = "complete" )
# Plot the obtained dendrogram
plot(hc1, cex = 0.6, hang = -1, main = "Dendrogram of Hierarchical Clustering")
rect.hclust(hc1, k = 10, border = 2:10)
```


Height

100% Natural Bran
Honey Comb
Golden Crisp
Smacks
Cocoa Puffs
Count Chocula
Froot Loops
Lucky Charms
Apple Jacks
Corn Flakes
Frosted Flakes
Fruity Pebbles
Golden Grahams
Nut & Honey Crunch
Wheaties Honey Gold
Double Chex
Cinnamon Toast Crunch
Cap'n Crunch
Honey Graham O's
Honey Nut Cheerios
Apple Cinnamon Cheerios
Crispy Wheat & Raisins

Quaker Oat Squares
Great Oats
Cracklin' Oat Bran
Raisin Nut Bran
Fruit & Raisin Bran
Post Nat.
Muesli Raisins
Dates & Almonds
Muesli Raisins & Walnuts
Fruit & Raisin Bran
Oatmeal Raisin Oats
Muesli Raisin Bran
Basic 4
Nutri-Grain Almond Raisin
Just Right Fruit & Nut
Total Raisin Bran
Special K
Special K
Rice Chex
Corn Flakes
Corn Flakes
Rice Krispies
Total Whole Grain
Total Crunchy
Nuggets
Just Right Total Corn Flakes
All Bran with Extra Fiber
All Bran
Purified White
Purified White
Raisin Squares
Frosted Mini-Wheats
Strawberry Fruit
Shredded Wheat
Shredded Wheat 'n Bran
Wheat Spoon Size
Nutri-Grain Flakes
Grape-Nuts
Multi-Grain Cheerios
Wheat Chex
Wheat Chex
Bran Flakes
Grape-Nuts

#

Computing with AGNES and with different linkage methods

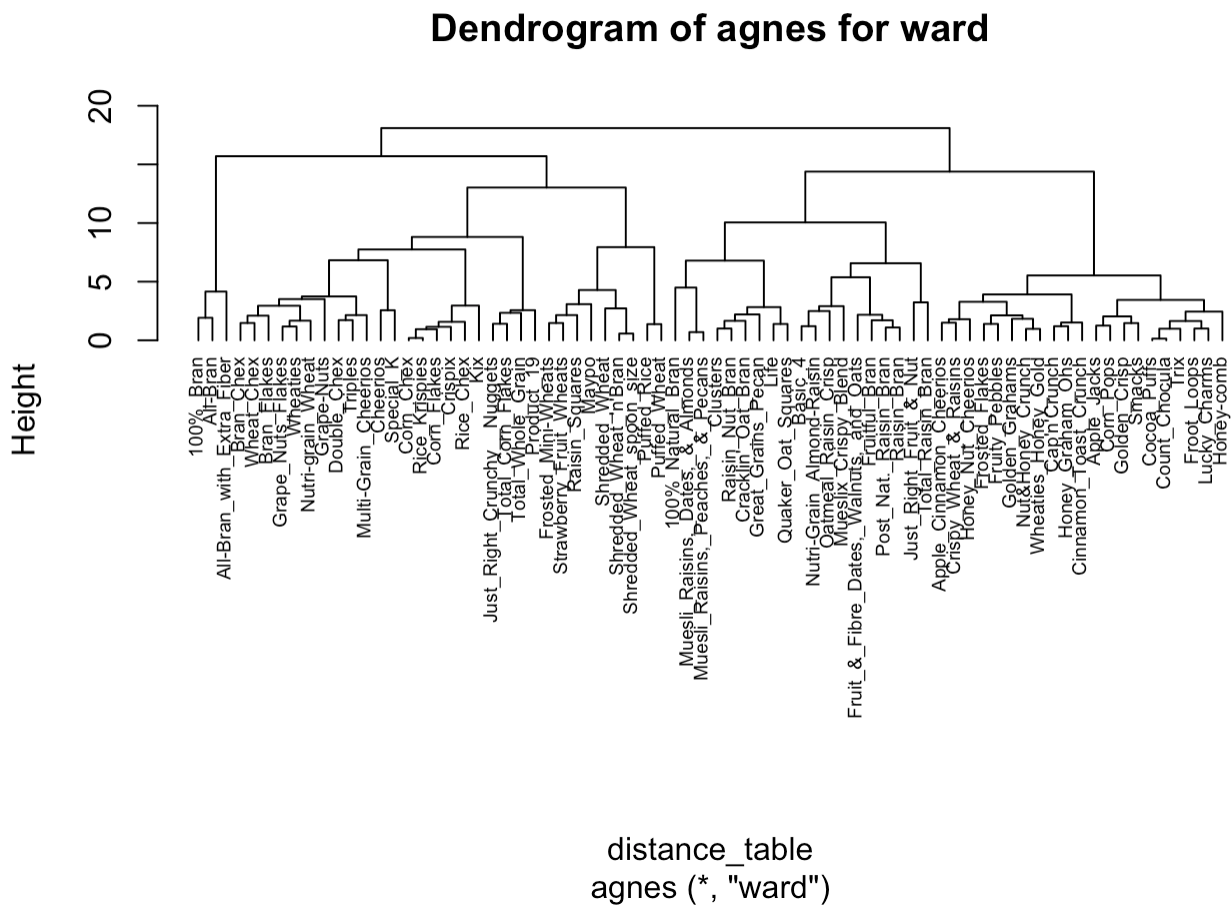
```
hc_ward <- agnes(distance_table, method = "ward")
print(hc_ward$ac)
```

```
## [1] 0.9087265
```

#These results confirm that the Ward linkage, which provides 90.87% accuracy, is the optimal agglomerative (AGNES) linkage to use.

Visualizing the dendrogram

```
hc_Ward <- agnes(distance_table, method = "ward")
pltree(hc_Ward, cex = 0.6, hang = -1, main = "Dendrogram of agnes for ward")
```



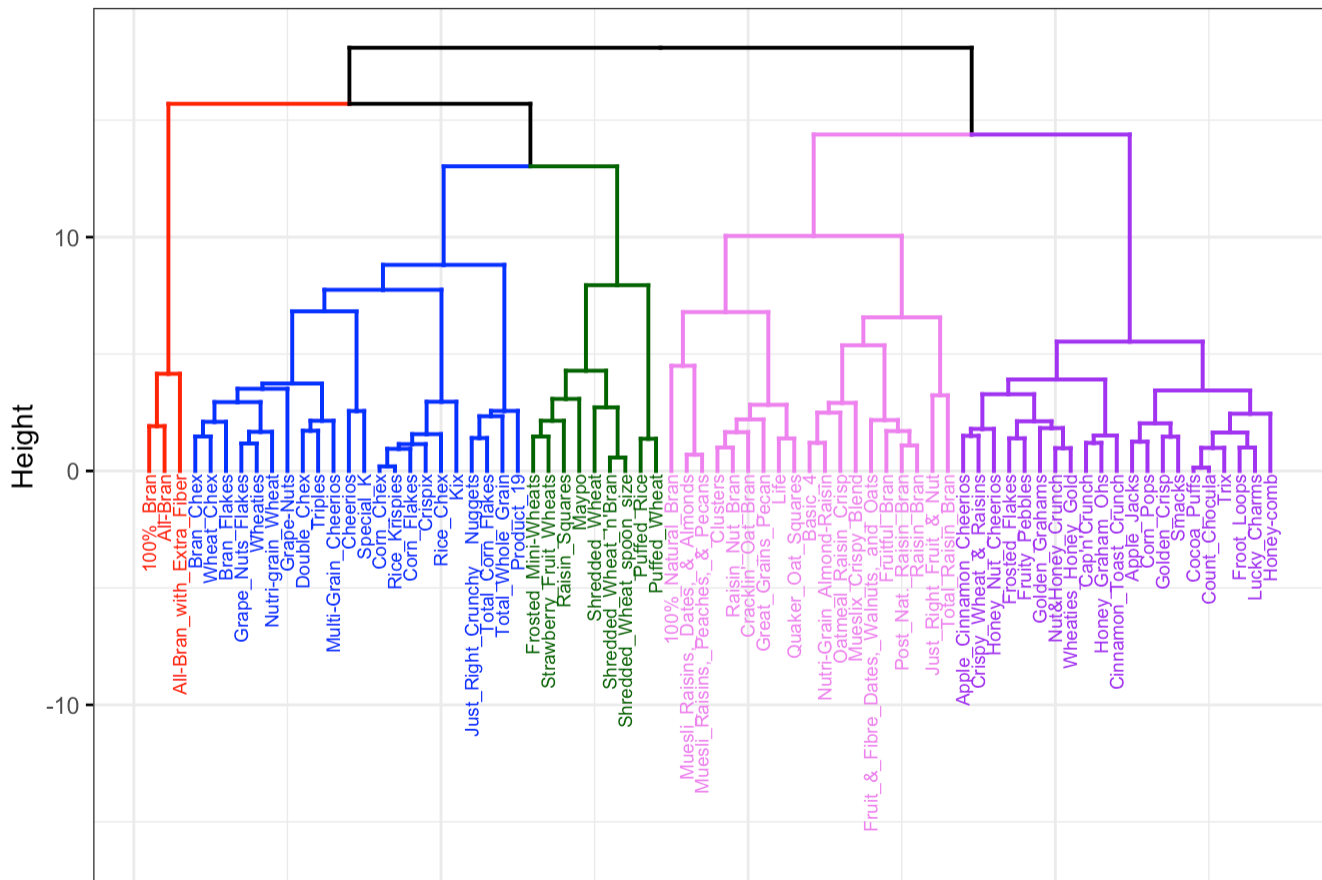
Question 2: How many cluster would you choose?

The largest difference in height can be used to determine the k value hence K =5 is the best option.

```
fviz_dend(hc_ward, k = 5, main = "Dendrogram of AGNES (Ward)", cex = 0.5, k_colors = c(
  "red", "blue", "darkgreen", "violet", "purple"), color_labels_by_k = TRUE, labels_track_height = 16, ggtheme = theme_bw())
```

```
## Warning: The `scale` argument of `guides()` cannot be `FALSE`. Use "none" instead as
## of ggplot2 3.3.4.
## i The deprecated feature was likely used in the factoextra package.
## Please report the issue at <]8;;https://github.com/kassambara/factoextra/issueshttps://github.com/kassambara/factoextra/issues[8;;>.
## This warning is displayed once every 8 hours.
## Call `lifecycle::last_lifecycle_warnings()` to see where this warning was
## generated.
```

Dendrogram of AGNES (Ward)



```
c_d6 <- cutree(hc_ward, k = 5)
Clustered_df <- as.data.frame(cbind ( c_d5, c_d6 ))
```

Question 3: Comment on the structure of the clusters and on their stability. Hint: To check stability, partition the data and see how well clusters formed based on one part apply to the other part

For the stability of the clusters, We will partition the data into A and B.

```
cereal_a = c_d5[1:55,]  
cereal_b = c_d5[56:74,]
```

Computing the distances of cereal_a

```
distance_cereal_a = get_dist(cereal_a)
```

#Compute with AGNES and with different linkage methods for cereal_a

```
hc_single_cereal_a <- agnes(distance_cereal_a, method = "single")  
print(hc_single_cereal_a$ac)
```

```
## [1] 0.6663587
```

```
hc_complete_cereal_a <- agnes(distance_cereal_a, method = "complete")  
print(hc_complete_cereal_a$ac)
```

```
## [1] 0.8285192
```

```
hc_average_cereal_a <- agnes(distance_cereal_a, method = "average")  
print(hc_average_cereal_a$ac)
```

```
## [1] 0.7646836
```

```
hc_ward_cereal_a <- agnes(distance_cereal_a, method = "ward")  
print(hc_ward_cereal_a$ac)
```

```
## [1] 0.8891086
```

With 88.91 % accuracy, it enables us to establish that the best linkage for cereal_a is Ward.

Computing the distances of cereal_a

```
distance_cereal_b = get_dist(cereal_b)
```

#Compute with AGNES and with different linkage methods for cereal_b

```
hc_single_cereal_b <- agnes(distance_cereal_b, method = "single")  
print(hc_single_cereal_b$ac)
```

```
## [1] 0.4805129
```

```
hc_complete_cereal_b <- agnes(distance_cereal_b, method = "complete")  
print(hc_complete_cereal_b$ac)
```

```
## [1] 0.71298
```

```
hc_average_cereal_b <- agnes(distance_cereal_b, method = "average")  
print(hc_average_cereal_b$ac)
```

```
## [1] 0.6232053
```

```
hc_ward_cereal_b <- agnes(distance_cereal_b, method = "ward")  
print(hc_ward_cereal_b$ac)
```

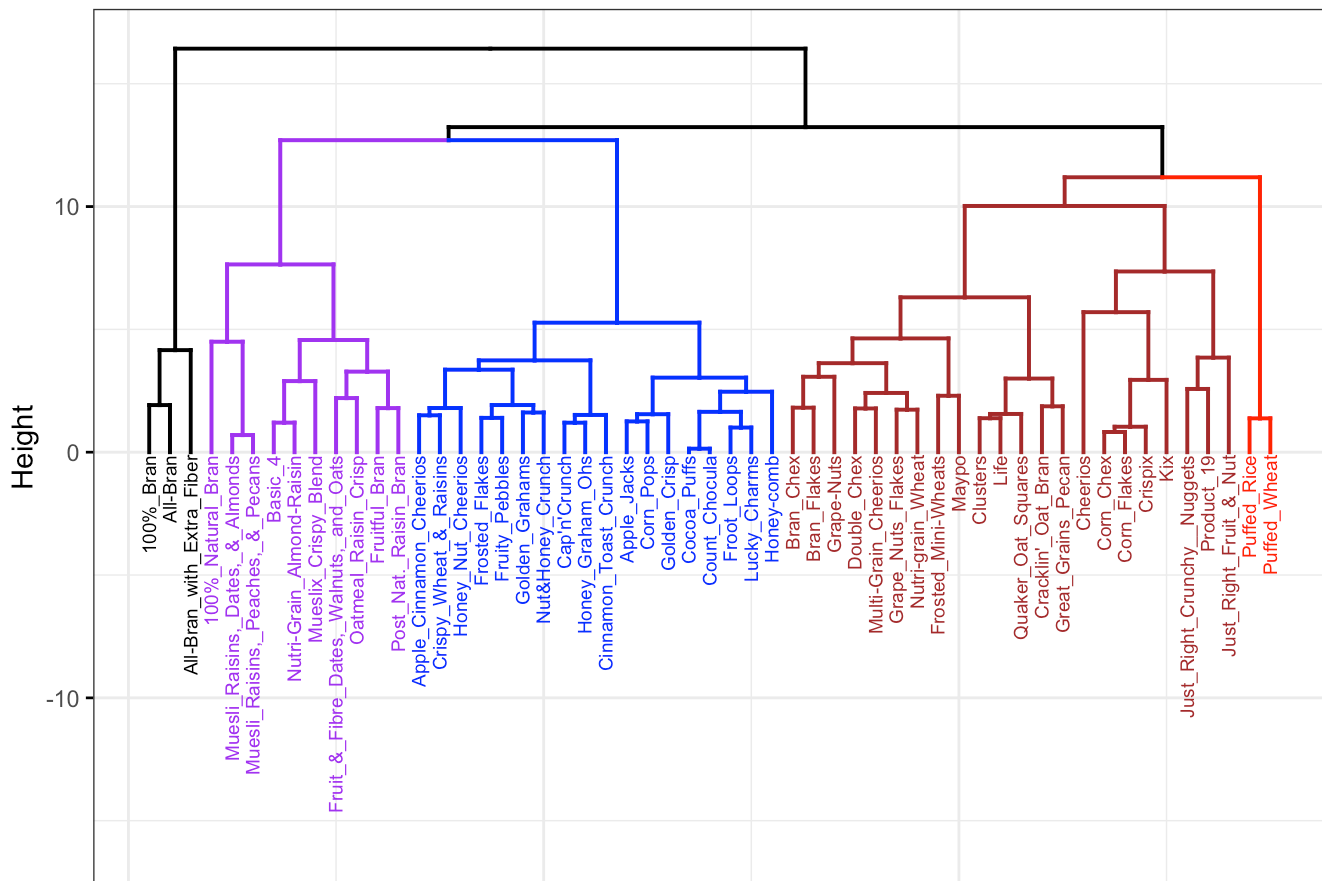
```
## [1] 0.7710122
```

With 77.10% accuracy, it enables us to establish that the best linkage for cereal_a is Ward.

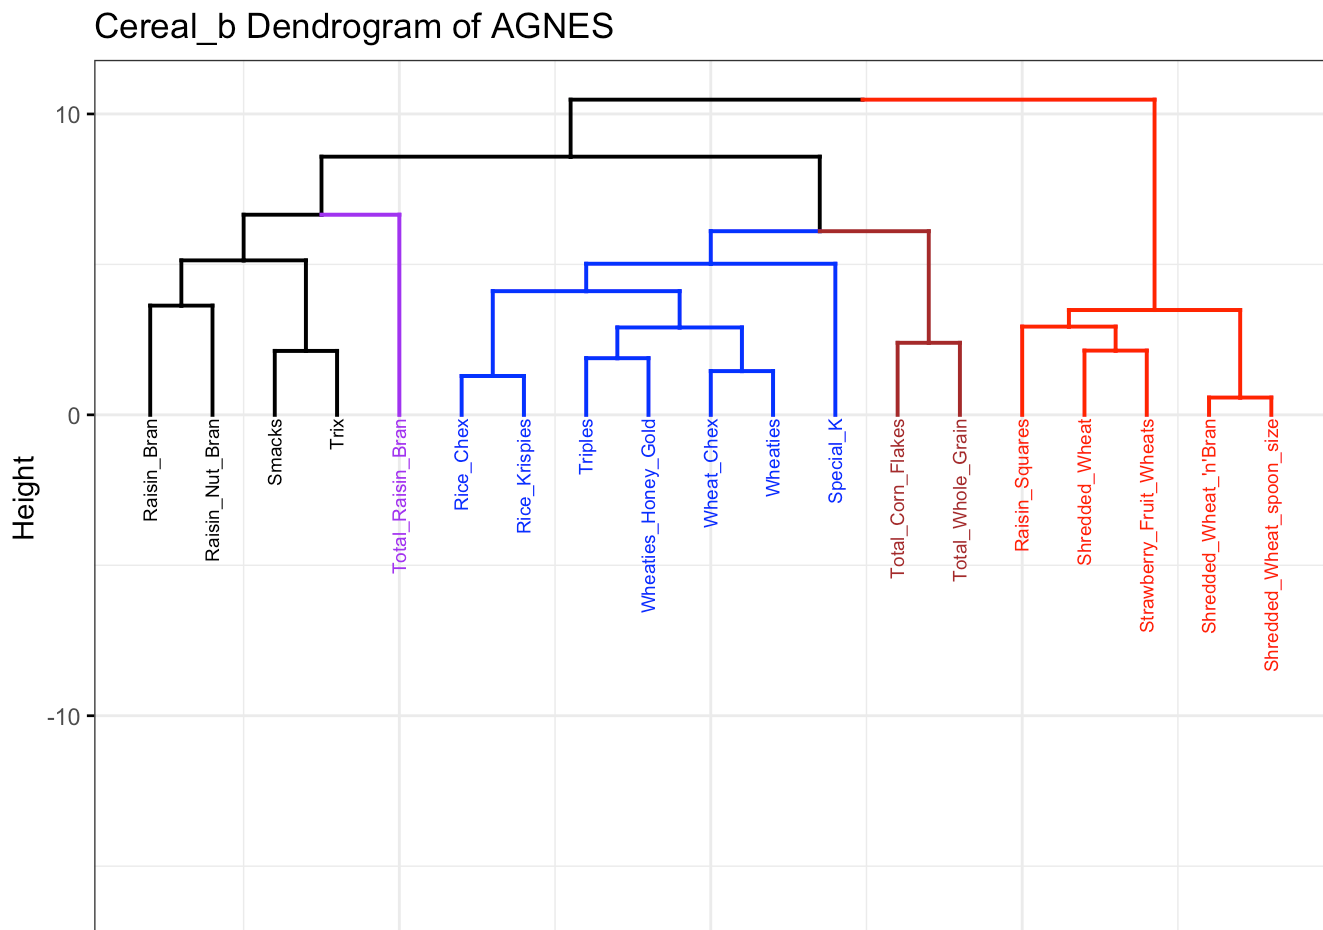
Plotting dendrogram of cereal_a and cereal_b

```
fviz_dend(hc_ward_cereal_a, k = 5, main = "Cereal_a Dendrogram of AGNES", cex = 0.5, k_  
colors = c("black", "purple", "blue", "brown", "red"), color_labels_by_k = TRUE, label_  
s_track_height = 16, ggtheme = theme_bw())
```

Cereal_a Dendrogram of AGNES



```
fviz_dend(hc_ward_cereal_b, k = 5, main = "Cereal_b Dendrogram of AGNES", cex = 0.5, k_
colors = c("black", "purple", "blue", "brown", "red"), color_labels_by_k = TRUE, label
s_track_height = 16, ggtheme = theme_bw())
```



Question 3 (part B): Use the cluster centroids from A to assign each record in partition B (each record is assigned to the cluster with the closest centroid)

```
Clustered_df_A <- cutree (hc_ward_cereal_a, k=5)
Clusters_A <- as.data.frame(cbind(cereal_a, Clustered_df_A))
Clust_1 <- colMeans (Clusters_A [Clusters_A$ Clustered_df_A == "1" ,])
# The centroid of cluster 1 is represented by a vector of mean values for each column
of the data as a result.
```

```
Clustered_df_B <- cutree (hc_ward_cereal_b, k=5)
Clusters_B <- as.data.frame(cbind(cereal_b, Clustered_df_B))
Clust_2 <- colMeans (Clusters_B [Clusters_B$ Clustered_df_B == "1" ,])
# The centroid of cluster 2 is represented by a vector of mean values for each column
of the data as a result.
```

```
Centroid <-rbind(Clust_1, Clust_2)
Centroid
```

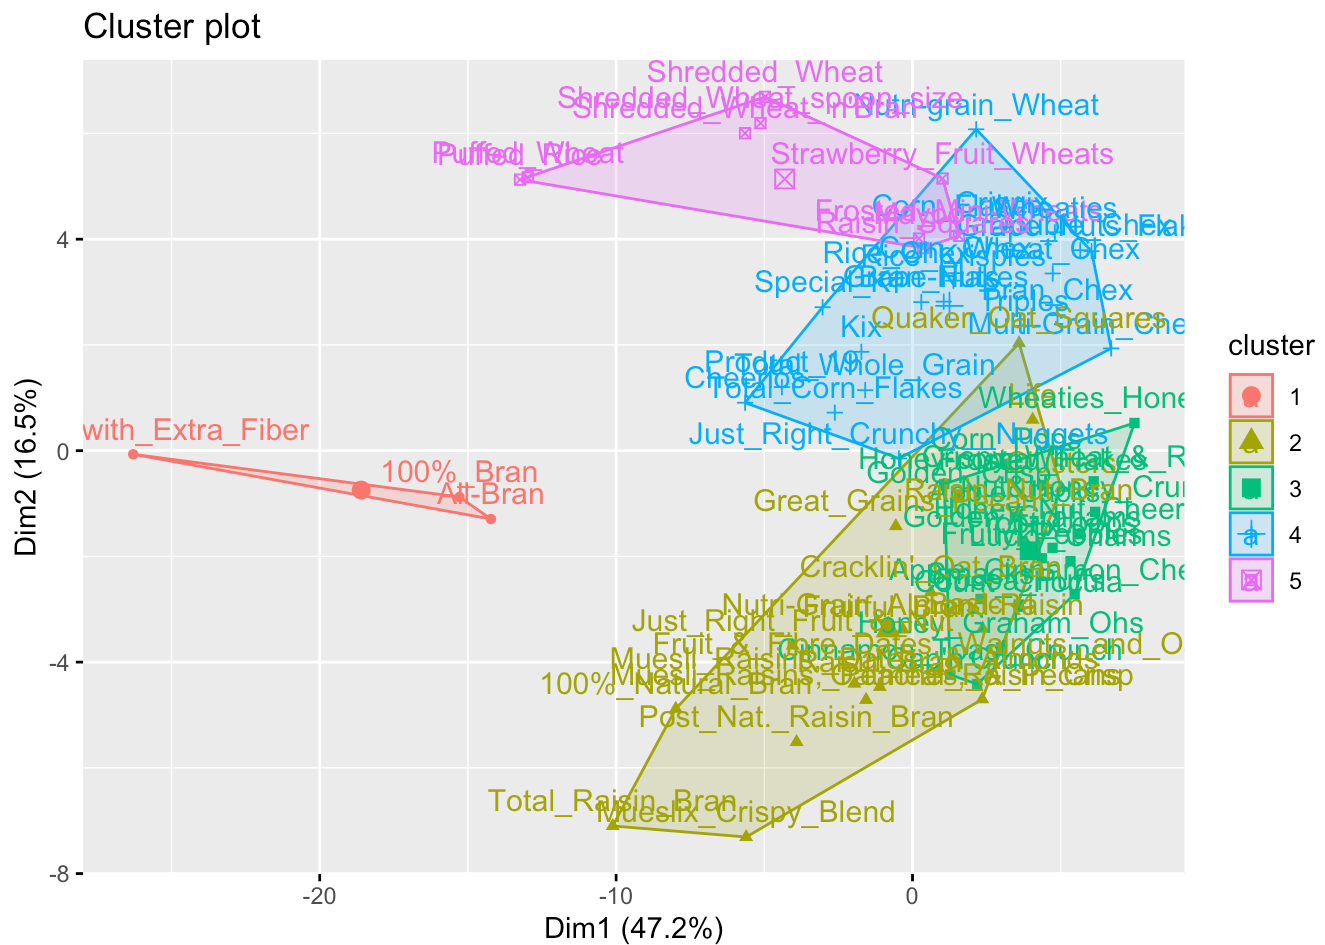
```
##          calories    protein      fat    sodium      fiber      carbo
## Clust_1 -2.201871  1.3817478 -0.3310734  0.1727901  3.64131237 -2.0718749
## Clust_2  0.149818 -0.2449462  0.2483051 -0.2702057 -0.02091106 -0.7977876
##          sugars    potass    vitamins    weight      cups      rating
## Clust_1 -0.7894824  2.9837813 -0.1818422 -0.2008324 -1.845255  2.2426479
## Clust_2  1.0648712  0.1796942 -0.1818422  0.3369228 -0.303848 -0.5618826
##          Clustered_df_A
## Clust_1                1
## Clust_2                1
```

Question 3 (part C): Assess how consistent the cluster assignments are compared to the assignments based on all the data.

After reviewing the centroid, it shows that cluster 1 is high in protein, fiber, and potassium. It means that the cereals in cluster 1 are more healthier than cluster 2. It can also be supported by looking at calories, fat, carbs, and sugar levels which are higher in cluster 2 as compared to cluster 1. Thus cereals in cluster 1 are healthier.

#Q4: The elementary public schools would like to choose a set of cereals to include in their daily cafeterias. Every day a different cereal is offered, but all cereals should support a healthy diet. For this goal, you are requested to find a cluster of “healthy cereals.” Should the data be normalized? If not, how should they be used in the cluster analysis?

```
#Visualizing the clusters in Scatter plot
fviz_cluster(list(data=distance_table, cluster = c_d6))
```

```
Healthy_cereal<- cbind(c_d2,c_d6)
mean(Healthy_cereal[Healthy_cereal$c_d6==1,"rating"])
```

```
## [1] 73.84446
```

```
mean(Healthy_cereal[Healthy_cereal$c_d6==2,"rating"])
```

```
## [1] 38.37137
```

```
mean(Healthy_cereal[Healthy_cereal$c_d6==3,"rating"])
```

```
## [1] 28.66112
```

```
mean(Healthy_cereal[Healthy_cereal$c_d6==4,"rating"])
```

```
## [1] 46.17608
```

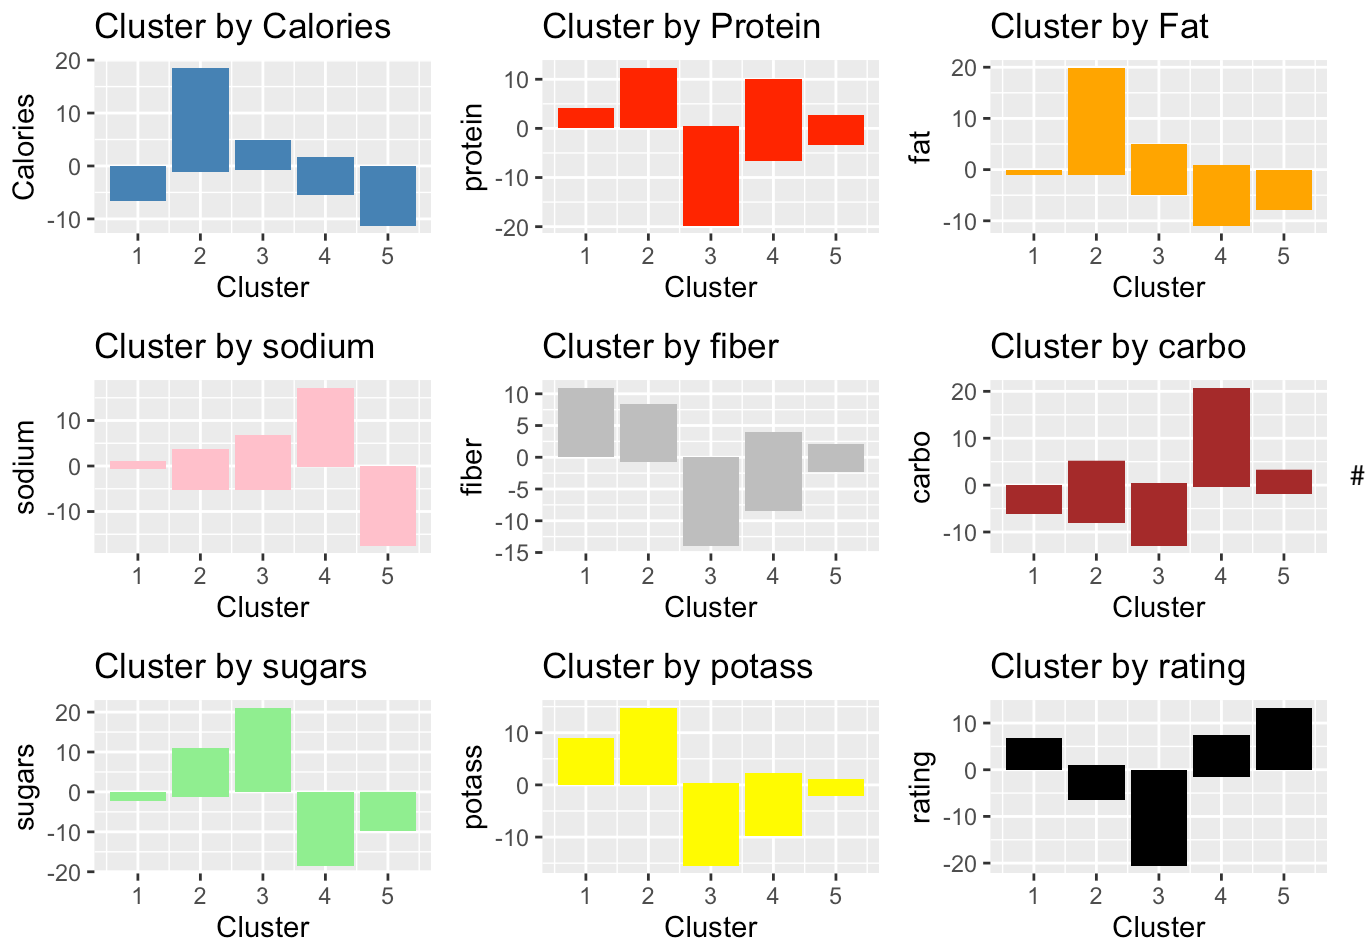
```
mean(Healthy_cereal[Healthy_cereal$c_d6==5,"rating"])
```

```
## [1] 63.0184
```

#It is evident that Cluster1 has the highest rating (73.84446), so we will select it as a nutritious cereal.

#lets also visualize the results by plotting a bar chart

```
calories <- ggplot(Clustered_df, aes(x = c_d6, y = calories)) +  
  geom_bar(stat = "identity", fill = "steelblue") +  
  labs(x = "Cluster", y = "Calories") +  
  ggtitle("Cluster by Calories")  
  
protein <- ggplot(Clustered_df, aes(x = c_d6, y = protein)) +  
  geom_bar(stat = "identity", fill = "red") +  
  labs(x = "Cluster", y = "protein") +  
  ggtitle("Cluster by Protein")  
  
fat <- ggplot(Clustered_df, aes(x = c_d6, y = fat)) +  
  geom_bar(stat = "identity", fill = "orange") +  
  labs(x = "Cluster", y = "fat") +  
  ggtitle("Cluster by Fat")  
  
sodium <- ggplot(Clustered_df, aes(x = c_d6, y = sodium)) +  
  geom_bar(stat = "identity", fill = "pink") +  
  labs(x = "Cluster", y = "sodium") +  
  ggtitle("Cluster by sodium")  
  
fiber <- ggplot(Clustered_df, aes(x = c_d6, y = fiber)) +  
  geom_bar(stat = "identity", fill = "gray") +  
  labs(x = "Cluster", y = "fiber") +  
  ggtitle("Cluster by fiber")  
  
carbo <- ggplot(Clustered_df, aes(x = c_d6,, y = carbo)) +  
  geom_bar(stat = "identity", fill = "brown") +  
  labs(x = "Cluster", y = "carbo") +  
  ggtitle("Cluster by carbo")  
  
sugars <- ggplot(Clustered_df, aes(x = c_d6,, y = sugars)) +  
  geom_bar(stat = "identity", fill = "lightgreen") +  
  labs(x = "Cluster", y = "sugars") +  
  ggtitle("Cluster by sugars")  
  
potass <- ggplot(Clustered_df, aes(x = c_d6,, y = potass)) +  
  geom_bar(stat = "identity", fill = "yellow") +  
  labs(x = "Cluster", y = "potass") +  
  ggtitle("Cluster by potass")  
  
rating <- ggplot(Clustered_df, aes(x = c_d6,, y = rating)) +  
  geom_bar(stat = "identity", fill = "black") +  
  labs(x = "Cluster", y = "rating") +  
  ggtitle("Cluster by rating")  
plot_grid(calories, protein, fat, sodium, fiber, carbo, sugars, potass, rating)
```



Here we can see that cluster 1 still has the best results. It is low in calories, sugar, and fat. It has higher content of fiber, potassium and protein. Thus we can conclude that cluster 1 can be a set of cereals to include in their daily cafeterias.